

Huanyi Xie

MS/PhD Student in Computer Science, KAUST | Advised by Prof. Di Wang
huanyi.xie@kaust.edu.sa | github.com/xiehuanyi | KAUST, Saudi Arabia | +966 54 941 9078

Efficient training and inference of large language models
LLM serving | Memory-efficient fine-tuning | Data attribution

Education

King Abdullah University of Science and Technology (KAUST) <i>MS/PhD in Computer Science</i>	Saudi Arabia Aug 2025 – Present
KTH Royal Institute of Technology <i>M.Sc. in Data Science</i>	Sweden Aug 2024 – Jun 2025
Harbin Institute of Technology, Weihai <i>B.Eng. in Artificial Intelligence</i>	China Sep 2020 – Jun 2024

Publications

TMLR'26	Efficient Text-Attributed Graph Learning through Selective Annotation and Graph Alignment Huanyi Xie, Lu Yu, Tianhao Huang, Longfei Li, Meng Li, Jun Zhou, Lijie Hu, Di Wang
ACL'26	High-Throughput and Memory-Efficient Zeroth-Order Fine-tuning LLMs with Distributed Parallel Computing Liangyu Wang, Huanyi Xie, Di Wang
ICLR'26	Predicting LLM Output Length via Entropy-Guided Representations Huanyi Xie, Yubin Chen, Liangyu Wang, Lijie Hu, Di Wang
ICLR'26	Dissecting Representation Misalignment in Contrastive Learning via Influence Function Huanyi Xie, Chenyang Ren, Khouloud Saadi, Shu Yang, Zhen Tan, Jingfeng Zhang, Lijie Hu, Di Wang
ICLR'26	Evaluating Data Influence in Meta Learning Chenyang Ren, Huanyi Xie, Shu Yang, Meng Ding, Dongrui Liu, Lijie Hu, Di Wang
COLM'25	ZO2: Scalable Zeroth-Order Fine-Tuning for Extremely Large LMs with Limited GPU Memory Liangyu Wang, Jie Ren, Hang Xu, Junxiao Wang, Huanyi Xie, David E. Keyes, Di Wang
ICCV'25	Semi-supervised Concept Bottleneck Models Lijie Hu, Tianhao Huang, Huanyi Xie, Xilin Gong, Chenyang Ren, Zhengyu Hu, Lu Yu, Ping Ma, Di Wang

Open-source Projects

nanoSGLang – Minimal LLM Serving Engine	https://github.com/xiehuanyi/nanoSGLang
- Built a high-performance LLM serving engine from scratch in Python/PyTorch, implementing continuous batching, chunked prefill, paged KV cache, prefix-sharing radix cache, streaming API, and request-level metrics. - Integrated FlashInfer paged attention and fused operators; benchmarked Qwen2.5-0.5B at 3469 tok/s on RTX A5000, exceeding the referenced SGLang baseline in the same workload.	
nanoMegatron – Minimal Distributed Training Framework	https://github.com/xiehuanyi/nanoMegatron
- Implemented core distributed-training algorithms in a compact educational codebase: DDP, ZeRO-1/2/3, tensor parallelism, sequence parallelism, pipeline parallelism, and expert parallelism. - Built a working SFT pipeline for Phi-tiny-MoE and GSM8K, with MoE AllToAll dispatch, ZeRO/FSDP-style sharding, and benchmark logs against DeepSpeed/PyTorch FSDP.	

Honors & Awards

Jun 2023	Honorable Mention, Mathematical Contest in Modeling (MCM)
Dec 2022	Gold Team Leader, Baidu PaddlePaddle Developer Groups
Jan 2022	Third Prize, Baidu Cognitive AI Creative Contest
Nov 2021	Second Prize, National College Student Mathematics Competition (Non-Mathematics Category)
Nov 2021	First Prize, China Collegiate Computing Contest Artificial Intelligence Innovation Contest